

rapidmesh

A Hierarchical Bottom-Up Centroidal-Voronoi Meshing Algorithm for B-Rep Solids with Analytic and NURBS Geometry

Specification and mathematical derivation

June 17, 2026

Abstract

This document specifies the meshing algorithm rapidmesh uses to turn a boundary representation (B-Rep) of one or more material regions into a conforming, isotropic, high-quality tetrahedral mesh. The algorithm is *dimensionally hierarchical*: it meshes 0-cells (corners), then 1-cells (edges), then 2-cells (faces), then the 3-cell (volume), and *freezes* each level before the next consumes it. Within every dimension it combines *error-driven adaptive sampling* (insert where a geometric tolerance is violated) with *sizing-weighted Centroidal Voronoi Tessellation* (Lloyd relaxation) for point quality. The frozen surface triangulation is a *hard constraint* on the volume Delaunay, which is what makes the boundary watertight by construction. We derive, with a computer-algebra system, the arc-length, curvature, deflection-sizing, first/second fundamental form and distance-faithful chart formulae for every edge and face modality the geometry kernel supports (line, circle/arc, surface-surface intersection, NURBS curve, extruded profile; plane, cylinder, cone, sphere, torus, extruded/ruled, NURBS), and give pseudocode for each stage.

Contents

1	Design philosophy	2
2	Notation and preliminaries	3
2.1	The sizing field	3
2.2	Centroidal Voronoi Tessellation	3
2.3	Discrete centroids and the restricted setting	3
2.4	Robust predicates	3
3	Stage 0 – Vertices	4
4	Stage 1 – Edges	4
4.1	Arc length	4
4.2	Curvature	4
4.3	Deflection sizing	4
4.4	1D CVT and error-driven adaptivity	5
5	Stage 2 – Faces	5
5.1	First fundamental form and area	6
5.2	Curvature and surface deflection sizing	6
5.3	Distance-faithful charts	6
5.4	Surface CVT and adaptivity	8

6	Stage 3 – Volume	8
6.1	Sizing field from the surface	8
6.2	Interior seeding	8
6.3	3D CVT and quality-driven adaptivity	8
6.4	Boundary-constrained tetrahedralization	8
7	The sizing field, end to end	9
8	Why this guarantees conformity	10

1 Design philosophy

Principle 1 (Dimensional hierarchy). A k -cell is meshed only after every $(k-1)$ -cell on its boundary is meshed and *frozen*. Corners pin edges; edges pin faces; faces pin the volume. Shared lower-dimensional entities are meshed once and reused, so adjacent cells agree on their common boundary exactly (conformity across seams is inherited, not reconstructed).

Principle 2 (Two orthogonal mechanisms per level). Each dimension uses (i) *error-driven adaptivity* to decide *how many* points and *where* (insert at the worst geometric-error location, then re-relax, until a tolerance is met), and (ii) *CVT/Lloyd relaxation* to decide *point quality* (move each free site to the density-weighted centroid of its (restricted) Voronoi cell). Adaptivity controls density; CVT controls regularity.

Principle 3 (Constrained volume, not recovered boundary). The frozen surface triangulation enters the volume stage as a set of *constraint facets*. The volume tetrahedralization is required to contain them as faces. The boundary is therefore *enforced*, not statistically recovered: there are no straddling tetrahedra, hence no “noses” protruding into a void and no gaps. This is the decisive structural property and the reason the surface must be completed *before* the volume is seeded.

Principle 4 (Embarrassingly parallel within each level). Freezing a level before the next consumes it makes the entities *within* a level mutually independent: they share only already-frozen, read-only lower-dimensional boundary. Hence each level is processed in parallel with no synchronisation:

- **all edges** at once (each depends only on its two fixed corners);
- then **all faces** at once (each depends only on its fixed boundary edges, which are computed once and read concurrently);
- then **all volume regions** at once (each material region’s constrained tetrahedralization depends only on its fixed bounding surface; shared interface facets are read-only).

Shared lower-dimensional entities are meshed exactly once and cached, which is what removes write contention. Each entity’s result is independently deterministic and merged in fixed id order, so parallelism does not affect reproducibility (§2). The three barriers (edges → faces → volume) are the only synchronisation points.

The geometry source is a B-Rep assembled from the exact CSG arrangement: faces are trimmed analytic (or NURBS) surfaces, edges are analytic curves (including the intersection curves booleans create), vertices are corner points, and the radial-edge topology records which faces meet along each edge and the material region on each side. The mesh is built *from this geometry*, independent of any input tessellation.

2 Notation and preliminaries

2.1 The sizing field

A scalar *sizing field* $h : \mathbb{R}^3 \rightarrow \mathbb{R}_{>0}$ prescribes the target element edge length. It is assembled bottom-up (edges set it on 1-cells, faces refine it on 2-cells, the volume smooths it through the interior) and is always made *Lipschitz* (gradation-limited):

$$h(\mathbf{x}) \leq h(\mathbf{y}) + \alpha \|\mathbf{x} - \mathbf{y}\| \quad \forall \mathbf{x}, \mathbf{y}, \quad (1)$$

with grading constant $\alpha \in (0, 1]$ (the default $\alpha = \frac{1}{2}$ grows neighbouring elements by at most $\sim 1.5\times$). Equation (1) is the standard mesh-gradation constraint [16]; it is enforced by a fast-marching / sweep over the size sources.

2.2 Centroidal Voronoi Tessellation

Given a density $\rho : \Omega \rightarrow \mathbb{R}_{>0}$ on a domain Ω and sites $\{\mathbf{z}_i\}_{i=1}^n$, the Voronoi region of $\mathbf{z}_i$ is $V_i = \{\mathbf{x} \in \Omega : \|\mathbf{x} - \mathbf{z}_i\| \leq \|\mathbf{x} - \mathbf{z}_j\| \ \forall j\}$. The tessellation is *centroidal* when each site coincides with the mass centroid of its region,

$$\mathbf{z}_i = \mathbf{c}_i := \frac{\int_{V_i} \rho(\mathbf{x}) \mathbf{x} \, d\mathbf{x}}{\int_{V_i} \rho(\mathbf{x}) \, d\mathbf{x}}. \quad (2)$$

A CVT is a critical point of the energy $\mathcal{F}(\{\mathbf{z}_i\}) = \sum_i \int_{V_i} \rho(\mathbf{x}) \|\mathbf{x} - \mathbf{z}_i\|^2 \, d\mathbf{x}$ [1]. *Lloyd's algorithm* [2] alternates “recompute V_i ” and “move $\mathbf{z}_i \leftarrow \mathbf{c}_i$ ” and converges to a CVT; the resulting point set is blue-noise / isotropic, which is precisely the input a Delaunay mesher needs for well-shaped elements.

Density that realises a sizing field. To make the local point spacing follow h , choose

$$\rho(\mathbf{x}) = h(\mathbf{x})^{-(d+2)} \quad (\text{theory, } d \text{ dimensions}), \quad \rho(\mathbf{x}) = h(\mathbf{x})^{-d} \quad (\text{used in practice}), \quad (3)$$

because a CVT equidistributes the energy density, giving local spacing $\propto \rho^{-1/(d+2)}$; the simpler h^{-d} (number density $\propto$ points per unit volume) is robust and is what we weight cells by. See [3, 4] for the exponent discussion.

2.3 Discrete centroids and the restricted setting

On a triangulation, the integral (2) is approximated by summing over the simplices incident to $\mathbf{z}_i$, each weighted by its measure times the density at its centroid. In dimension d with incident simplices T of centroid $\mathbf{g}_T$, measure $|T|$,

$$\mathbf{c}_i \approx \frac{\sum_{T \ni \mathbf{z}_i} |T| \rho(\mathbf{g}_T) \mathbf{g}_T}{\sum_{T \ni \mathbf{z}_i} |T| \rho(\mathbf{g}_T)}. \quad (4)$$

On a curved surface, the relevant object is the *restricted* Voronoi diagram (Voronoi cells clipped to the surface) and its dual *restricted Delaunay* triangulation [6, 7]; in practice we relax in a 2D chart (§5) so (4) applies in the chart and the move is lifted back onto the surface.

2.4 Robust predicates

All combinatorial decisions (orientation, in-circle/in-sphere) use exact adaptive-precision predicates [14]; only non-decision quantities (centroid weights, sizing values) are floating point. Map/iteration order is made deterministic (seedless hashing), so the mesh is reproducible.

3 Stage 0 – Vertices

Every B-Rep corner $\mathbf{p}$ becomes a *pinned* site: it never moves and is present in every downstream triangulation. Pinning corners preserves sharp features exactly. Formally these are the 0-cells of the cell complex and the fixed points of all relaxations below.

4 Stage 1 – Edges

A B-Rep edge is a curve $\mathbf{C} : [t_0, t_1] \rightarrow \mathbb{R}^3$ between two corners. We support the modalities in Table 1.

Modality	$\mathbf{C}(t)$	Source
Line	$\mathbf{p}_0 + t \mathbf{d}$	straight B-Rep edge
Circle / arc	$\mathbf{c} + r(\cos t \mathbf{x} + \sin t \mathbf{y})$	rim, fillet, boolean of quadrics
Intersection	implicit $\mathbf{S}_a \cap \mathbf{S}_b$, seeded by the chain	generic surface–surface boolean
NURBS	$\frac{\sum_i N_{i,p}(t) w_i \mathbf{P}_i}{\sum_i N_{i,p}(t) w_i}$	imported / profile curve [17]
Extruded profile	$\mathbf{b} + z \mathbf{a} + u_x P_x(t) + u_y P_y(t)$	lifted 2D profile (airfoil)
Polyline	piecewise linear chain	faceted fallback

Table 1: Edge curve modalities. All expose $\mathbf{C}(t), \mathbf{C}'(t), \mathbf{C}''(t)$.

4.1 Arc length

The arc length and the natural (unit-speed) parameter are

$$s(t) = \int_{t_0}^t \|\mathbf{C}'(\tau)\| d\tau, \quad ds = \|\mathbf{C}'(t)\| dt. \quad (5)$$

We tabulate $s(t)$ once (adaptive Gauss/Simpson) and invert $t(s)$ by table lookup + a Newton step using $dt/ds = 1/\|\mathbf{C}'\|$, so points can be placed at prescribed arc-length spacings. For a circle $\|\mathbf{C}'\| = r$ and $s = r \Delta t$ exactly; for a line it is exact; for NURBS / extruded profiles the table is used.

4.2 Curvature

For a space curve the (unsigned) curvature is

$$\kappa(t) = \frac{\|\mathbf{C}'(t) \times \mathbf{C}''(t)\|}{\|\mathbf{C}'(t)\|^3}, \quad R(t) = \frac{1}{\kappa(t)} \text{ (radius of curvature)}, \quad (6)$$

derived from the Frenet apparatus and verified symbolically. A straight segment has $\kappa \equiv 0$, $R \equiv \infty$; a circle of radius r has $\kappa \equiv 1/r$.

4.3 Deflection sizing

The chord between two samples a parametric distance apart deviates from the true arc by the *sagitta*. For an arc of radius R subtending a half-angle $\theta/2$, with chord length $\ell = 2R \sin(\theta/2)$, the maximum chord–arc deviation is

$$\varepsilon = R(1 - \cos \frac{\theta}{2}) = \frac{R \theta^2}{8} + O(\theta^4) \approx \frac{\ell^2}{8R}. \quad (7)$$

Solving the small-angle relation for the admissible edge length given a *deflection tolerance* ε gives the local target

$$\boxed{\ell(\mathbf{x}) = \sqrt{8\varepsilon R(\mathbf{x})}} \quad \left(\text{exact: } \ell = 2\sqrt{R^2 - (R - \varepsilon)^2}\right), \quad (8)$$

both verified with sympy. The edge sizing field is then

$$h_{\text{edge}}(\mathbf{x}) = \min\left(h_{\max}, \sqrt{8\varepsilon R(\mathbf{x})}\right), \quad (9)$$

i.e. refine where the curve bends (R small), coast at $h_{\max}$ where it is straight. This is the curve analogue of the surface bound (13).

4.4 1D CVT and error-driven adaptivity

On a fixed edge the n interior sites $s_1 < \dots < s_n$ (arc-length coordinates, endpoints pinned) relax by the 1D centroid rule (the segment analogue of (4)): with neighbours s_{i-1}, s_{i+1} and density $\rho = h^{-1}$,

$$s_i \leftarrow \frac{\rho_{i-\frac{1}{2}} \frac{s_{i-1} + s_i}{2} + \rho_{i+\frac{1}{2}} \frac{s_i + s_{i+1}}{2}}{\rho_{i-\frac{1}{2}} + \rho_{i+\frac{1}{2}}}, \quad \rho_{i \pm \frac{1}{2}} = h\left(\frac{s_i + s_{i \pm 1}}{2}\right)^{-1}, \quad (10)$$

which equidistributes spacing $\propto h$. Adaptivity wraps the relaxation:

Algorithm 1 MESHEDGE($\mathbf{C}, h, \varepsilon$)

```

1: sample  $s(t)$  table;  $L \leftarrow s(t_1)$ 
2:  $n \leftarrow \max(1, \lceil \int_0^L h(s)^{-1} ds \rceil)$  ▷ initial count from the sizing field
3: place  $s_1, \dots, s_n$  equidistributing  $h$  (cumulative- $\rho$  inversion)
4: repeat
5:   for a few iterations do relax all  $s_i$  by (10)
6: end for
7:  $e^*, s^* \leftarrow$  worst chord-arc deflection and its location ▷ error-driven
8:  $\ell^*, m^* \leftarrow$  longest segment length and its midpoint ▷  $h_{\max}$ -driven
9: if  $e^* > \varepsilon$  or  $\ell^* > h_{\max}$  then insert a site (at  $s^*$  if  $e^* > \varepsilon$ , else  $m^*$ );  $n += 1$ 
10: end if
11: until  $e^* \leq \varepsilon$  and  $\ell^* \leq h_{\max}$  and max site move  $< \tau_{\text{move}}$ 
12: return  $\{\mathbf{C}(s_i)\}$  ▷ frozen edge points; corners included

```

The relax-insert loop is the 1D instance of the general scheme: relax for quality, insert at the worst error, repeat until the tolerance holds. Shared edges are meshed once and cached, so both incident faces read identical points.

5 Stage 2 – Faces

A face is a trimmed surface $\mathbf{S} : \Omega \subset \mathbb{R}^2 \rightarrow \mathbb{R}^3$, $(u, v) \mapsto \mathbf{S}(u, v)$, bounded by oriented loops of frozen edges. The strategy is: parameterise the face by a *distance-faithful chart*, relax interior points there with a 2D surface-CVT (boundary = frozen edge points, held fixed), insert where the surface deflection tolerance is violated, and lift every chart point exactly onto $\mathbf{S}$. The output is a *conforming triangulation of the face* that the volume stage will treat as a hard constraint.

Face categories and the face sizing field. Three chart families cover all modalities (§5.3): *planar* faces use the plane chart; *developable* analytic faces (cylinder, cone, extruded) use an exact isometric unroll; *doubly-curved* analytic faces (sphere, torus) and *periodic* full revolutions use a

distance-faithful chart or the periodic revolution construction, with a restricted-CVT fallback for general NURBS. Before any face is seeded, its sizing field is assembled from the *now-frozen edge sampling* on its boundary (the edge spacings and any per-region $h_{\max}$) combined with the surface deflection bound (13), then Lipschitz-smoothed (1) across the face. The frozen edges therefore both bound the face and seed its density.

5.1 First fundamental form and area

With $\mathbf{S}_u = \partial \mathbf{S} / \partial u$, $\mathbf{S}_v = \partial \mathbf{S} / \partial v$, the metric (first fundamental form) coefficients and area element are

$$E = \mathbf{S}_u \cdot \mathbf{S}_u, \quad F = \mathbf{S}_u \cdot \mathbf{S}_v, \quad G = \mathbf{S}_v \cdot \mathbf{S}_v, \quad dA = \sqrt{EG - F^2} \, du \, dv. \quad (11)$$

A curve $t \mapsto (u(t), v(t))$ in parameter space has length $\int \sqrt{E\dot{u}^2 + 2F\dot{u}\dot{v} + G\dot{v}^2} \, dt$; this is the metric the chart must reproduce for Euclidean 2D operations (separation, centroid) to mean arc length on the surface. Table 2 lists E, F, G, dA for every analytic modality, all derived with sympy.

Surface	$\mathbf{S}(u, v)$	E	F	G	dA
Plane	$\mathbf{o} + u\mathbf{e}_1 + v\mathbf{e}_2$	1	0	1	$du \, dv$
Cylinder (θ, h)	$r(\cos \theta, \sin \theta) + h\mathbf{a}$	r^2	0	1	$r \, du \, dv$
Sphere (λ, φ)	$r(\cos \varphi \cos \lambda, \cos \varphi \sin \lambda, \sin \varphi)$	$r^2 \cos^2 \varphi$	0	r^2	$r^2 \cos \varphi \, du \, dv$
Cone (θ, s)	$s \tan \alpha (\cos \theta, \sin \theta) + s\mathbf{a}$	$s^2 \tan^2 \alpha$	0	$\sec^2 \alpha$	$ s \tan \alpha \sec \alpha \, du \, dv$
Torus (θ, ϕ)	$(R+r \cos \phi)(\cos \theta, \sin \theta) + r \sin \phi \mathbf{a}$	$(R+r \cos \phi)^2$	0	r^2	$r R+r \cos \phi \, du \, dv$
Extruded (t, h)	$\mathbf{b} + P(t) + h\mathbf{a}$	$\ P'(t)\ ^2$	0	1	$\ P'(t)\ \, du \, dv$
NURBS (u, v)	tensor B-spline [17]	$\mathbf{S}_u \cdot \mathbf{S}_u$	$\mathbf{S}_u \cdot \mathbf{S}_v$	$\mathbf{S}_v \cdot \mathbf{S}_v$	$\sqrt{EG - F^2}$

Table 2: First fundamental form per modality ($\mathbf{a}$ the unit axis; α the cone half-angle). All but the NURBS row are diagonal ($F = 0$), which is what makes the charts below clean.

5.2 Curvature and surface deflection sizing

The shape operator’s eigenvalues are the principal curvatures κ_1, κ_2 ; from the second fundamental form $L = \mathbf{S}_{uu} \cdot \mathbf{n}$, $M = \mathbf{S}_{uv} \cdot \mathbf{n}$, $N = \mathbf{S}_{vv} \cdot \mathbf{n}$ ($\mathbf{n} = \mathbf{S}_u \times \mathbf{S}_v / \|\mathbf{S}_u \times \mathbf{S}_v\|$), the Gaussian and mean curvatures are

$$K = \kappa_1 \kappa_2 = \frac{LN - M^2}{EG - F^2}, \quad H = \frac{1}{2}(\kappa_1 + \kappa_2) = \frac{EN - 2FM + GL}{2(EG - F^2)}, \quad (12)$$

and $\kappa_{1,2} = H \pm \sqrt{H^2 - K}$. Verified values (Table 3) confirm the geometric intuition: developables have $K = 0$, the sphere is umbilic.

The surface sagitta argument mirrors (7) with $R = R_{\min} = 1/\kappa_{\max}$: a triangle edge of length ℓ on a surface of smallest principal radius $R_{\min}$ deviates from the surface by $\varepsilon \approx \ell^2/(8R_{\min})$, so the face sizing field is

$$h_{\text{face}}(\mathbf{x}) = \min\left(h_{\max}, \sqrt{8\varepsilon R_{\min}(\mathbf{x})}\right), \quad R_{\min} = 1/\kappa_{\max}. \quad (13)$$

5.3 Distance-faithful charts

A chart $\Phi : \text{face} \rightarrow \mathbb{R}^2$ lets the planar CVT machinery act on the surface. We want Φ as close to an *isometry* as possible so Euclidean 2D distance equals surface arc length.

Surface	K	principal curvatures	$\kappa_{\max} = 1/R_{\min}$
Plane	0	0, 0	0
Cylinder	0	$1/r$, 0	$1/r$
Cone	0	$1/(s \tan \alpha \sec \alpha)$, 0	$1/(s \tan \alpha \sec \alpha)$
Sphere	$1/r^2$	$1/r$, $1/r$	$1/r$
Torus	$\frac{\cos \phi}{r(R + r \cos \phi)}$	$\frac{1}{r}$, $\frac{\cos \phi}{R + r \cos \phi}$	$1/r$
Extruded	0	$\kappa_P(t)$, 0	$\kappa_P(t) = \frac{ P' \times P'' }{\ P'\ ^3}$

Table 3: Principal/Gaussian curvature per modality (sympy-verified). $R_{\min} = 1/\kappa_{\max}$ feeds the surface deflection bound.

Developables (cylinder, cone, extruded): exact isometry. A surface with $K \equiv 0$ unrolls into the plane without distortion.

- *Cylinder*: $\Phi(\theta, h) = (r\theta, h)$. Then $E = G = 1, F = 0$ in the chart – exact.
- *Cone*: polar unroll $\Phi(\theta, s) = (\rho \cos \psi, \rho \sin \psi)$ with slant $\rho = s/\cos \alpha$ and sweep $\psi = \theta \sin \alpha$. The induced chart metric is $E = s^2 \tan^2 \alpha$, $F = 0$, $G = \sec^2 \alpha$ – identical to the cone’s first fundamental form (Table 2), i.e. *exactly isometric* (sympy-verified).
- *Extruded*: $\Phi(t, h) = (\sigma(t), h)$ with $\sigma(t) = \int^t \|P'\|$ the profile arc length; $E = G = 1, F = 0$ – exact.

Sphere: azimuthal-equidistant. The sphere ($K = 1/r^2 \neq 0$) admits no isometry to the plane. We use the azimuthal-equidistant chart about a cap centre: the chart radial coordinate is the exact geodesic distance $r\psi$ (ψ the geodesic colatitude), so radial distances are exact; the circumferential direction stretches by

$$\text{stretch}(\psi) = \frac{\psi}{\sin \psi} = 1 + \frac{\psi^2}{6} + \frac{7\psi^4}{360} + O(\psi^6), \quad (14)$$

(sympy-verified) which is ≤ 1.05 for caps up to $\psi \approx 32^\circ$ and stays bounded well past a hemisphere. The chart is a bijection on any cap that excludes the antipode; bijectivity is validated by a round-trip test $\text{dist}(\text{proj}_{\mathbf{S}}(\mathbf{p}), \Phi^{-1}\Phi(\mathbf{p})) < \text{tol}$.

General NURBS / non-developable: metric-scaled or restricted CVT. For a NURBS patch we scale the parameters by the local metric $(\sqrt{E}, \sqrt{G})$ at the patch centroid (a first-order isometry) and, when a single chart distorts too much, fall back to a genuine restricted-Voronoi relaxation on the surface [5, 4]: the centroid (4) is computed from the restricted cells and the move is projected back onto $\mathbf{S}$ by closest-point.

Revolution faces and periodicity. A full revolution (cylinder/cone/torus barrel) is periodic in θ ; its boundary loops are rim circles that do not close in the unrolled rectangle. Such a face is meshed by a structured (θ, v) grid whose rings reuse the *rim* longitudes (radial alignment, no seam) and whose row spacing follows h in arc length – the uniform grid is the on-surface CVT optimum of a developable. A sphere cap meeting an intersection circle is itself a revolution about the line of centres (the circle lies in a plane $\perp$ that axis), so the same latitude/longitude construction applies, with both incident caps deriving their longitude reference from the *shared* rim so their rings align.

5.4 Surface CVT and adaptivity

Algorithm 2 MESHFACE($\mathbf{S}, \text{loops}, h, \varepsilon$)

```

1: build chart  $\Phi$  (Table-driven by modality); validate bijectivity
2:  $B \leftarrow \Phi(\text{frozen edge points of all loops})$   $\triangleright$  fixed boundary in chart
3: scatter interior sites in  $\Phi(\text{face})$  at spacing  $h$  (greedy graded Poisson) inside the trim loops
4: repeat
5:   for a few iterations do
6:     2D Delaunay of  $B \cup \{\text{interior}\}$  (exact predicates)
7:     move each interior site by (4),  $\rho = h^{-2}$ , clamp inside the loops; keep  $B$  fixed
8:   end for
9:    $e^* \leftarrow \max \text{triangle deflection vs. } \mathbf{S}$  (sagitta, (13))  $\triangleright$  error-driven
10:   $\ell^* \leftarrow \max \text{triangle edge length (in arc length)}$   $\triangleright h_{\max}$ -driven
11:  if  $e^* > \varepsilon$  or  $\ell^* > h_{\max}$  then insert a site at the offending triangle's circumcentre (in the chart)
12:  end if
13: until  $e^* \leq \varepsilon$  and  $\ell^* \leq h_{\max}$  and  $\max \text{move} < \tau_{\text{move}}$ 
14: lift:  $\mathbf{p}_k = \mathbf{S}(\Phi^{-1}(\text{chart site}_k))$  – exact on surface
15: return conforming triangulation of the face (its triangles + vertices)

```

After Stage 2 every face carries a frozen, conforming triangulation; shared edges guarantee face-to-face conformity. The union of all face triangulations is a watertight surface mesh $\mathcal{S}$.

6 Stage 3 – Volume

6.1 Sizing field from the surface

The volume sizing field is grown from the completed surface: at each surface vertex the local feature size and the surface element size seed h_{vol} , which is then propagated through the interior and gradation-limited by (1). Crucially this is computed *after* the surface is final, so it already reflects curvature- and feature-driven refinement, including *local thickness*: where two boundary sheets are closer than $h_{\max}$ (a thin web, a blunt trailing edge, a slender bore wall) the field shrinks so the interior is resolved.

6.2 Interior seeding

Interior sites are thrown at the field density: a randomised, separation-guarded dart pattern (or a body-centred-cubic lattice graded by h [15]) with a per-point clearance from $\mathcal{S}$ of one local element, and a global budget $\int_{\Omega} h^{-3} dV$ so a sharp feature does not inflate the count. Surface vertices are *not* re-seeded; they are the boundary of the volume point set.

6.3 3D CVT and quality-driven adaptivity

Interior sites relax by the 3D centroid rule (4) with $\rho = h^{-3}$; surface vertices stay fixed. A site whose centroid leaves the domain is mirrored back across the nearest boundary plane. After relaxation, sites are inserted where the local edge length exceeds $h_{\max}$ or where element quality (minimum dihedral angle, radius–edge ratio) is poor, and the relaxation is repeated.

6.4 Boundary-constrained tetrahedralization

This is the structural heart of the method. The volume is the *constrained Delaunay tetrahedralization* (CDT) of the interior sites *plus* the frozen surface mesh $\mathcal{S}$ as constraint facets

[13, 12, 10]:

$$\text{Tetrahedralize}(\{\text{interior sites}\} \cup V(\mathcal{S})) \quad \text{s.t. every } \triangle \in \mathcal{S} \text{ is a union of tet faces.} \quad (15)$$

Because $\mathcal{S}$ is watertight and every constraint triangle is recovered as tetrahedron faces, each surface triangle has a tetrahedron exactly on each side; region labels are then assigned by flood fill across non-constraint faces starting from the region tag attached to each surface triangle's two sides. No tetrahedron straddles the boundary, so:

Per-region parallelism. Each material region is bounded by a disjoint subset of the frozen $\mathcal{S}$ (its outer walls and the interface facets it shares with neighbours, all read-only), so the regions' constrained tetrahedralizations are independent and run in parallel (Principle 4); a shared interface facet is meshed once on $\mathcal{S}$ and both regions tetrahedralize against the same fixed triangles, which keeps the multi-material mesh conforming across interfaces.

Proposition 1 (Watertightness). *If $\mathcal{S}$ is a closed, non-self-intersecting, conforming surface mesh and the CDT recovers all of its facets, then for every constraint triangle the two incident tetrahedra carry the region labels prescribed by that triangle's front/back tags, and the boundary of each meshed region is exactly the corresponding subset of $\mathcal{S}$. In particular there are no protrusions into a void and no gaps.*

Facet recovery in a CDT may require a bounded number of Steiner points *on the constraint facets*; these are inserted on $\mathcal{S}$ (never off it), so the boundary geometry is preserved exactly [12, 13]. Sliver tetrahedra (near-zero volume, good circumradius) that survive Delaunay refinement are removed by a final quality pass – vertex smoothing constrained to carriers, plus flips/exudation [11, 9] – which is volume- and conformity-preserving because boundary vertices move only within their surface/edge carrier.

Algorithm 3 MESHVOLUME($\mathcal{S}, h$)

- 1: build h_{vol} from $\mathcal{S}$ (feature size, thickness), gradation-limit by (1)
 - 2: seed interior sites at density h_{vol}^{-3} , clearance $\geq h$ from $\mathcal{S}$
 - 3: **repeat**
 - 4: CDT of (interior $\cup V(\mathcal{S})$) with $\mathcal{S}$ as constraints
 - 5: **for** a few iterations **do** relax interior sites by (4), $\rho = h^{-3}$
 - 6: **end for**
 - 7: insert sites where edge length $> h_{\text{max}}$ or quality is poor
 - 8: **until** size and quality targets met **and** max move $< \tau_{\text{move}}$
 - 9: flood-fill region labels across non-constraint faces from surface tags
 - 10: quality post-pass (carrier-constrained smoothing, flips); drop slivers
 - 11: **return** tetrahedra, per-tet region, boundary faces ($= \mathcal{S}$)
-

7 The sizing field, end to end

The single field h threads all stages and is only ever *refined* downward in size:

1. **Edges** set h from curve deflection (9) and h_{max} / per-region overrides.
2. **Faces** take the min of the edge field and the surface deflection bound (13); the smoothed result drives the surface CVT.
3. **Volume** takes the min of the face field and the local-thickness / feature-size estimate, then gradation-limits (1) through the interior.

Because each stage only ever lowers h and then re-smooths, the field is monotone and Lipschitz, and the element-size contract (a documented $1.5\times$ max-edge bound, as in gmsh) holds globally.

8 Why this guarantees conformity

The earlier rapidmesh path built one *unconstrained* Delaunay over all points and recovered the boundary by labelling each tetrahedron with the region of its centroid. That recovers the boundary only *statistically*: where the surface sample is coarser than a feature, a tetrahedron straddles the true surface, and whole-tet labelling then produces a protrusion (centroid inside) or a gap (centroid outside). The present algorithm removes the failure mode at its root by (a) completing a watertight, error-controlled surface mesh *first*, and (b) making that mesh a *constraint* on the volume rather than a target to be rediscovered. Conformity is then a theorem (Proposition 1), not a sampling-density gamble. The cost is the constrained Delaunay / facet recovery, which is exactly the machinery the method must keep; the variational (CVT) relaxation supplies the well-shaped interior that keeps recovery cheap and slivers rare [8, 9].

References

- [1] Q. Du, V. Faber, M. Gunzburger. *Centroidal Voronoi Tessellations: Applications and Algorithms*. SIAM Review 41(4), 1999.
- [2] S. P. Lloyd. *Least squares quantization in PCM*. IEEE Trans. Information Theory 28(2), 1982.
- [3] Q. Du, D. Wang. *Anisotropic centroidal Voronoi tessellations and their applications*. SIAM J. Sci. Comput. 26(3), 2005.
- [4] B. Lévy, Y. Liu. L_p Centroidal Voronoi Tessellation and its applications. ACM TOG (SIGGRAPH) 29(4), 2010.
- [5] D.-M. Yan, B. Lévy, Y. Liu, F. Sun, W. Wang. *Isotropic Remeshing with Fast and Exact Computation of Restricted Voronoi Diagram*. Computer Graphics Forum (SGP) 28(5), 2009.
- [6] H. Edelsbrunner, N. R. Shah. *Triangulating topological spaces*. Int. J. Comp. Geom. Appl., 1997.
- [7] J.-D. Boissonnat, S. Oudot. *Provably good sampling and meshing of surfaces*. Graphical Models 67, 2005.
- [8] P. Alliez, D. Cohen-Steiner, M. Yvinec, M. Desbrun. *Variational Tetrahedral Meshing*. ACM TOG (SIGGRAPH) 24(3), 2005.
- [9] J. Tournois, C. Wormser, P. Alliez, M. Desbrun. *Interleaving Delaunay Refinement and Optimization for Practical Isotropic Tetrahedron Mesh Generation*. ACM TOG (SIGGRAPH) 28(3), 2009.
- [10] S.-W. Cheng, T. K. Dey, J. R. Shewchuk. *Delaunay Mesh Generation*. CRC Press, 2012.
- [11] S.-W. Cheng, T. K. Dey, H. Edelsbrunner, M. Facello, S.-H. Teng. *Sliver Exudation*. J. ACM 47(5), 2000.
- [12] J. R. Shewchuk. *Constrained Delaunay Tetrahedralizations and Provably Good Boundary Recovery*. 11th Int. Meshing Roundtable, 2002.
- [13] H. Si. *TetGen, a Delaunay-Based Quality Tetrahedral Mesh Generator*. ACM TOMS 41(2), 2015.
- [14] J. R. Shewchuk. *Adaptive Precision Floating-Point Arithmetic and Fast Robust Geometric Predicates*. Discrete & Comp. Geom. 18, 1997.

- [15] F. Labelle, J. R. Shewchuk. *Isosurface Stuffing: Fast Tetrahedral Meshes with Good Dihedral Angles*. ACM TOG (SIGGRAPH) 26(3), 2007.
- [16] H. Borouchaki, F. Hecht, P. J. Frey. *Mesh gradation control*. Int. J. Numer. Meth. Eng. 43, 1998.
- [17] L. Piegl, W. Tiller. *The NURBS Book*. 2nd ed., Springer, 1997.
- [18] M. P. do Carmo. *Differential Geometry of Curves and Surfaces*. Prentice-Hall, 1976.
- [19] P. J. Frey, P.-L. George. *Mesh Generation: Application to Finite Elements*. 2nd ed., Wiley, 2008.